

REFocus: *Real-world motion-Blur Recovery via a luminance-First conditional Colorization pipeline* — A Two-Stage Approach for Deblurring and Optional Colorization

Nikhitha Nagalla
Applied Data Science
University of Florida
Gainesville, Florida - 32608
nikhithanagalla@ufl.edu

Abstract— REFocus is a two-stage image restoration system that first reconstructs structure in the luminance domain and then colorizes only when needed. Stage-1 employs DeblurGAN to deblur and inpaint on the L channel, while Stage-2 uses a CycleGAN with ResNet-block generators to map grayscale/faded content to plausible color without paired supervision. Since Deliverable-2, we standardized Lab preprocessing and range clamping, introduced a warm-start schedule (content/SSIM before adversarial loss) for DeblurGAN, and increased identity loss for CycleGAN to reduce hue drift. We also refined the web UI into a clear two-module flow with per-stage previews and downloads. On a 300-image test set, colorization achieved MSE = 0.0054, PSNR = 24.32 dB, and SSIM = 0.9180; DeblurGAN’s test content loss plateaued at ≈ 0.114 , indicating stable luminance restoration. Loss curves reveal emerging overfitting in Stage-2 after ~ 25 – 30 epochs, motivating stronger regularization and broader palette diversity. Planned work includes ΔE (Lab) reporting, inside- vs outside-mask metrics, overlap-add tiling for very high-resolution inputs, and optional user color hints with an edit-heatmap for transparency.

Keywords— deblurring, inpainting, colorization, DeblurGAN, CycleGAN, ResNet generators, Lab color space, perceptual loss, web interface.

I. INTRODUCTION

1.1 Problem Context and Project Summary

Smartphones, action cams, and surveillance systems frequently capture motion-blurred, defocused, or noisy scenes, especially people walking, vehicles in motion, and low-light street shots. Family snapshots also suffer from handshake blur and poor exposure. One-click “enhance” filters often oversharpen or introduce artifacts, while single-stage models struggle to handle real camera blur.

Refocus is a two-stage pipeline: Stage 1 restores structure on luminance (reduces blur/noise,

recovers edges). Stage 2 colorizes only when the image is grayscale (e.g., mono CCTV or scanned B&W photos), preserving original luminance/texture. Delivered as a drag-and-drop web tool with Before/After views, Refocus aims to produce clear, natural images that are easier to review, search, and share, helping both families and analysts.

II. DATASET

For the two phases, I want to use two datasets. One for restoring the images and one for coloring the images.

2.1. HIDE (Human-Aware Motion Deblurring) from Kaggle - Restoration

Type & content: Paired RGB images with realistic human-centric motion blur (walking people, handheld cameras, street scenes) and corresponding sharp ground truth captured with short exposure. Pairs reflect non-uniform, human-induced blur, closer to surveillance and everyday snapshots than synthetic kernels.

Intended use: Train Stage-1 Restorer in Lab space on L channel only (supervision: $L_{\text{blur}} \rightarrow L_{\text{sharp}}$). Decoupling structure (L) from color avoids hue shifts and lets the same model serve both grayscale and color inputs.

Scale & splits: The dataset will be around 16GB with thousands of pairs. We use the official train/val/test lists (or generate a consistent split) and patch-based training at 256×256 .




Fig. 1. Dataset

2.2 COCO 2017 (train2017 + val2017), via Hugging Face - Colorization

Type & content: Large, diverse RGB image collection spanning people, objects, and scenes; ideal for learning broad color priors. We treat it as unpaired data for $L \rightarrow ab$ prediction.

Intended use: Train Stage-2 Colorizer to predict chrominance (ab) from luminance (L). In inference, run colorization only if the restored image is effectively grayscale. Otherwise, preserve original color ($ab_{original}$).

Scale & splits: ~118k train2017 images and 5k val2017 images (~18 GB total if fully mirrored). We create an internal validation split from train2017 (~95/5) and use val2017 as the test set for final reporting.

III. PREPROCESSING

Stage 1: Restoration (HIDE)

- Each training example has two images: a blurry photo and its sharp version.
- Then, converted both to Lab color and used only the L channel (lightness/brightness).
- scaled L to 0–1 and clip values so nothing is outside the range.
- Train: From the input image, then center-crop to 256×256 .
- Apply the same crop to blurry and sharp, and sometimes flip left/right.
- Val/Test: just center-crop to 256×256 .
- From HIDE’s training images, we keep about 90% for training and 10% for validation.
- We use HIDE’s test set as the final check.

Stage 2: Colorization (COCO via Hugging Face)

- We download images on the fly from official COCO URLs (and cache them on disk).

- Convert each image to Lab. Input to the model: L (lightness) in 0–1
- Target to learn: ab (color) in -1 to 1
- Clipped after conversion to keep values in range.
- Train: From the input, center-crop to 256×256 , sometimes flip.
- Val/Test: center-crop to 256×256 , no flip.
- COCO train2017 (~112k images) $\rightarrow$ we make our own split:
- 95% train (106,655 images) and 5% validation (5,613 images).
- COCO val2017 (5,000 images) is our test set.
- Many images have little color. To stop the model from becoming “too gray,” we give more weight to pixels that actually have strong color when computing the loss.

Ethical / Privacy considerations

state how long uploads are kept (or don’t store them at all); provide a “Delete after download” option. use images only for research; document the source; avoid faces in public demos without consent.

IV. PLANNED ARCHITECTURE




Fig.2. Architecture

1. Input will be Image RGB
2. Preprocess (convert to brightness + color layers, resize/crop, clamp)
3. Stage-1: DeblurGAN (Restoration on brightness only)
 - Generator restores the brightness of the image (blurry $\rightarrow$ sharp)
 - Discriminator checks the realism of the restored brightness
 - Output: restored brightness image

“Is it grayscale?” check

- If the original has little/no color → go to Stage-2
- Else keep the original color layers

4. Stage-2: CycleGAN (Colorization)

- Domain X: grayscale (brightness duplicated to 3 channels)
- Domain Y: natural color photos (from COCO)
- Generator X→Y: gray → color (predicts color layers or full RGB)
- Generator Y→X: color → gray
- Cycle + identity losses keep hues stable and realistic
- Output: predicted color layers (or a color image you can convert)

5. Recompose: combine restored brightness + chosen color layers → RGB → Streamlit UI showing both the blurred and sharpened images.

4.1 Updated System Architecture

Pipeline (v2):

Data → Lab conversion and crops → Stage-1 DeblurGAN on L → “Is-grayscale?” gate → (if gray) Stage-2 CycleGAN → Recompose (L + ab) → UI.

Preprocessing: Standardized Lab conversion with range clamping; consistent 256×256 center-crops for train/val/test; color-aware weighting for pixels with strong chroma.

Model/HPs: Warm-start DeblurGAN with content/SSIM before adversarial loss; tuned identity loss for CycleGAN to preserve hue when input is already colored; maintained PatchGAN discriminators.

Interface flow: Mode toggles (“Restore only”, “Restore+Colorize”, “Always colorize”), auto gray detection, and side-by-side preview with planned download/export. (See earlier architecture + UI figures.)

V. USER INTERFACE PLAN

I want to build the interface in **Streamlit** because it supports rapid prototyping, easy side-by-side image display, and simple deployment. The app will run locally during development and can be hosted on Streamlit Community Cloud for demos. The interface will accept a single image upload in JPG, PNG, or TIFF format. In the sidebar, users will

choose a processing mode like “Restore only,” “Restore + Colorize,” or “Always colorize”. They can toggle “Auto-detect grayscale” to run colorization only when the input has little or no color. The app will display a **side-by-side** comparison. The original image on the left and the processed image on the right, each scaled to fit the page and usable with the standard Streamlit zoom behavior.




Fig. 3. Sample UI

VI. INNOVATION AND ANTICIPATED CHALLENGES

Domain gap (real scenes vs. training data).

Risk: Low light, compression, or fine artifacts may be under-represented.

Mitigation: Add input-only augmentations (vignetting, JPEG noise, light scratches), run early checks on a small curated set, and, if needed, fine-tune on a few real examples.

GAN stability and over/under-restoration.

Risk: Adversarial training can ring or hallucinate. Pure content loss can oversmooth.

Mitigation: Warm-start DeblurGAN with content/SSIM before enabling GAN loss. Use PatchGAN with instance/spectral norm. Monitor PSNR/SSIM and visual grids. Expose a “deblur strength” slider to blend results in the UI.

Desaturation or hue drift in colorization.

Risk: The model may produce dull colors or incorrect tints.

Mitigation: Use a color-aware (chroma-weighted) loss, keep identity loss in CycleGAN, and always replace the output’s brightness with the restored brightness before rendering. Provide a “restore-only” option and trigger colorization only when grayscale is detected.

VII. IMPLEMENTATION TIMELINE

Week	Focus	Example
Oct 20 - 26	Set up HIDE/COCO loaders, preprocessing, fixed splits; scaffold DeblurGAN (restore) + CycleGAN (color).	Data ready; baseline scripts run.
Oct 27 – Nov 2	Train DeblurGAN (warm-start with content/SSIM, then GAN); add val PSNR/SSIM.	Stable restorer + checkpoints.
Nov 3 – 16	Train CycleGAN (gray↔color) with identity + cycle + color-aware loss; compose restored brightness + color.	Plausible colorization; compose working.
Nov 17 – 30	Build Streamlit UI: upload, modes, auto-gray, strength slider, face-blur, EXIF strip, side-by-side view. Refinement	Working interactive demo. Refinement
Dec 1 – 11	Rehearse demo. Finalize report, GitHub, and slides.	Presentation-ready deliverables.

Table.1. Timeline

VIII. MODELS IMPLEMENTED

Frameworks and Libraries

Python [3.10/3.11], PyTorch [2.x] (torchvision), OpenCV, NumPy, scikit-image, matplotlib; Streamlit or Gradio for the demo UI. Notebooks used: image-restoration.ipynb (restoration) and Colorization.ipynb (colorization).

Training Setup.

- **Frameworks:** PyTorch (AMP), OpenCV/scikit-image.
- **Hardware:** [GPU]. **Epochs/Batch:** DeblurGAN [50 ep, B=8]; CycleGAN [2 ep, B=4].

- **Optimizers:** AdamW (lr 2e-4, betas 0.5, 0.999).
- **Augmentations:** Random crops/flips, synthetic scratch/stain masks, JPEG artifacts.

DeblurGAN (Restoration).

Generator: U-Net/ResNet-style deblurring generator with residual blocks and skip connections.

Discriminator: PatchGAN.

Losses: Adversarial + content (L1/Charbonnier on luminance) + perceptual (VGG) + TV near mask edges.

Use: Prioritized loss inside damage masks; outputs a structurally restored image

CycleGAN (Colorization) — ResNet Generators.

Architecture: Two **ResNet-block generators** ($X \rightarrow Y$ and $Y \rightarrow X$) and two PatchGAN discriminators. Each generator stacks an initial conv + N residual blocks (e.g., 6 for 256px, 9 for 512px) + upsampling.

Domains: X = grayscale/faded artworks; Y = clean/color artworks.

Losses: Adversarial (for each direction) + **cycle-consistency** ($\|G(F(y)) - y\|_1 + \|F(G(x)) - x\|_1$) + **identity loss** (stabilizes hue when input already colored).

Use: Takes DeblurGAN output and predicts chroma while keeping edges/texture from the restored luminance.

IX. EVALUATION RESULTS

The Preliminary evaluation for the models cycleGAN and DeblurGAN are as follows. This chart shows smoothed test metrics per batch for the colorization stage: the purple line (ΔE_{mean} , MA5) stays roughly in the 16–18 range across batches, indicating stable perceptual color error. The green (MAE_RGB, MA5) and orange (MAE_ab, MA5) lines remain low and flat, suggesting no batch-specific degradation during evaluation.




Fig.4. Preliminary Result - CycleGAN

Regarding the test metrics of DeblurGAN, I got a loss of 0.1154.

On the image test set, the colorization stage reports MSE = 0.0054, PSNR = 24.32 dB, and SSIM = 0.9180. The training curve for colorization shows steady train-loss decline from ≈ 0.0052 to ≈ 0.0041 , while val-loss fluctuates and trends slightly upward ($\approx 0.0052 \rightarrow \approx 0.00545$), indicating emerging overfitting after ~ 25 – 30 epochs.




Fig. 5. Training vs validation loss - Colorization(U-Net)

For Stage-1 restoration, DeblurGAN-v2 remains stable: train G-loss hovers around 0.096–0.099, val content-loss stays near 0.098–0.101, and test content-loss plateaus at ≈ 0.114 across 50 epochs.




Fig. 6. Preliminary results 1

Interpretation. Stage-1 achieves consistent luminance restoration with a flat validation/test content-loss curve, suggesting convergence but limited headroom without architectural or loss changes. Stage-2 produces perceptually strong

results (SSIM ≈ 0.92) at modest PSNR, but the widening train–val gap signals overfitting; regularization and data diversity should help.




Fig. 7. Loss curves comparison - Image Restoration

What changed vs. Deliverable 2. (i) More consistent Lab preprocessing and warm-start scheduling stabilized DeblurGAN, holding test content-loss near ~ 0.114 (previous runs varied around ~ 0.115). (ii) Identity/cycle tuning improved hue stability, reflected in higher SSIM, though PSNR remains typical for translation models.

X. REFINEMENTS MADE SINCE DELIVERABLE 2

- Cleaner preprocessing: Uniform Lab pipeline; clipped ranges ($L \in [0, 1]$, $ab \in [-1, 1]$); consistent crops; synchronized flips across pairs.
- Training stability: DeblurGAN warm-start (content-first), then enable GAN; instance/spectral norm on PatchGAN; mixed precision kept on.
- Hyperparameters: AdamW ($2e-4$, $\beta=0.5/0.999$), TV near mask edges; tuned CycleGAN identity/cycle weights to reduce hue shift.
- Updated loops: Checkpointing with embedded configs; seed control; simple LR scheduling; validation every N steps with L-channel PSNR/SSIM.
- Interface/UX: Auto grayscale gate, clearer mode labels, and progress planning; export (PNG/TIFF) with note “AI-assisted restoration” (in progress).
- Error handling: Input type checks, EXIF strip option, and size guards (tiling planned for very large images).

XI. USER INTERFACE

The REFocus UI is a single-page, dark-theme web app with a centered title (“Refocus”) and tagline that explains the two-stage flow. The webpage features two vertically stacked modules, where Stage 1 (Deblur/Restore) is the first one and Stage 2 (Colorization/Final Output) is the second. Each stage has an input card on the left and an output card on the right. During Stage 1, the users can drag-and-drop a JPG/PNG file, click on the Run Stage 1 button, and then a huge restored preview will appear along with the option to Download Output; in Stage 2, they can upload a new image or click on Use Stage 1 Output, and then Run Stage 2 to colorize and download the final result.




Fig. 8. User Interface - Image Restoration

This arrangement provides a straightforward flow—Upload → Run Stage 1 → (optional) Use Stage 1 Output → Run Stage 2 → Download—while the mirrored cards, inline format hints, and per-stage downloads make it easier to navigate. Unlike Deliverable 2, the explicit handoff button and clearer microcopy have jointly improved usability thus fastened the process of side-by-side verification. The remaining enhancement consists of a progress indicator for large images and an optional before/after slider.




Fig. 9. User Interface - Image Colorization

Interface Usability and Improvements

What changed: Centered hero text (“REFocus”), concise intro, left/right image cards with filenames,

and a simple control strip (mode toggle + “Show results”). Usability notes: Immediate filename feedback; deterministic behavior for “Restore only”; conditional colorization avoids unnecessary style shifts for already-colored images. (See UI plan and sample UI in prior report.)

XII. CHALLENGES AND NEXT STEPS

Challenges

- Data realism: Few examples of authentic damage; synthetic masks miss crackle, fray, and stains.
- Palette variance: Limited diversity across art styles causes color drift in CycleGAN.
- Large defects: DeblurGAN struggles with very big missing regions.
- GAN stability: Sensitive to loss weights; occasional over-sharpened textures.
- Compute limits: High-res inputs need tiling; naive tiling creates seams.
- Evaluation gap: PSNR/SSIM rise, but perceptual quality varies; few true color references.
- UX: Slow inference on large images; no user color hints or edit-region visualization.

Next Steps

- Data & masks: Add a small real-damage validation set; upgrade synthetic masks (crackle, stains, frayed tears).
- Loss tuning: Edge-aware + light frequency loss for DeblurGAN; higher identity loss and soft palette prior for CycleGAN.
- Training strategy: Small hyperparameter sweep; curriculum from small → large defects; keep AMP and add EMA/R1 for stability.
- Inference quality: Overlap-add tiling with blending; optional multi-scale refinement; gentle tone mapping.
- Evaluation: Report inside-mask vs. outside-mask metrics; add LPIPS and ΔE ; run a small A/B human rating.
- Interface polish: Progress bar, batch mode, export with “AI-assisted restoration” note; optional user color hint and edit heatmap.
- Reproducibility: Finalize [requirements.txt](#), [config.yaml](#), and versioned checkpoints with embedded configs.

XIII. RESPONSIBLE AI REFLECTION

Fairness & Representation. Colorization can shift skin tones or cultural artifacts. I will include diverse examples in validation, review failures, and keep conservative defaults. A “Restore only” mode lets users avoid AI-generated color when it’s inappropriate. Privacy & Consent. Surveillance-like images may contain bystanders. The app will offer face-blur and EXIF stripping by default in demos. I will avoid redistributing input images and process files in memory or delete them after download. Data Licensing & Provenance. I will respect COCO/HIDE licenses and keep links/attribution. I will not re-share raw datasets; I will share only derived results where allowed and record the dataset/source in experiment logs.

REFERENCES

- [1] O. Kupyn, V. Budzan, M. Mykhailych, D. Mishkin, and J. Matas, “DeblurGAN: Blind motion deblurring using conditional adversarial networks,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2018, pp. 8183–8192.
- [2] O. Kupyn, T. Martyniuk, J. Wu, and Z. Wang, “DeblurGAN-v2: Deblurring (orders-of-magnitude) faster and better,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2019, pp. 8878–8887.
- [3] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros, “Unpaired image-to-image translation using cycle-consistent adversarial networks,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 2242–2251.
- [4] R. Zhang, P. Isola, and A. A. Efros, “Colorful image colorization,” in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2016, pp. 649–666.
- [5] P. Isola, J.-Y. Zhu, T. Zhou, and A. A. Efros, “Image-to-image translation with conditional adversarial networks,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 1125–1134. (PatchGAN)
- [6] W. Shen, R. Zhang, Y. Xu, and J. Ma, “Human-aware motion deblurring,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2019, pp. 5571–5580. (HIDE dataset)